Jeremiah Ochepo
CSC 8015 Data Mining & Predictive Analytics Spring 2024
Assignment #7 Data Streaming
Principal Components Analysis - PCA
Due Date: April 12, 2024

Introduction:

In this analysis, we explore the characteristics and relationships within the "mtcars" dataset using various statistical and machine-learning techniques. The dataset comprises observations on different car models, with measurements such as miles per gallon (mpg), number of cylinders (cyl), displacement (disp), horsepower (hp), and other specifications. We aim to extract insights into the factors influencing mileage, identify distinct car types, and build predictive models to understand the dataset more comprehensively.

Defining Variables:

- **mpg:** Miles per gallon, representing fuel efficiency.
- **cyl:** Number of cylinders in the engine.
- **disp:** Displacement, the total volume of all cylinders in the engine.
- **hp:** Horsepower, a measure of engine power.
- **drat:** Rear axle ratio.
- **wt:** Weight of the car.
- **qsec:** Quarter mile time (seconds).
- **vs:** Engine type (0 for V-shaped, 1 for straight).
- **am:** Transmission type (0 for automatic, 1 for manual).
- **gear:** Number of forward gears.
- **carb:** Number of carburetors.

NaN Values and Dimension Incompatibility:

We meticulously checked for NaN (Not a Number) values throughout the analysis to ensure data integrity. Fortunately, no NaN values were detected, indicating the dataset's cleanliness and completeness. However, during the matrix multiplication process, we encountered an error indicating incompatible dimensions. This error suggests that the matrices involved in the multiplication operation have mismatched dimensions, hindering the computation. Further investigation into the dimensionality of the matrices and potential corrective actions will be necessary to proceed with the analysis effectively.

Answers to Questions:

1. **Should a principal components analysis of this data be based on the covariance or the correlation matrix? Explain.**

   A principal components analysis (PCA) of this data should be based on the correlation matrix rather than the covariance matrix. Since the variables in the dataset have different scales, using the correlation matrix ensures that each variable contributes equally to the analysis, regardless of its scale.

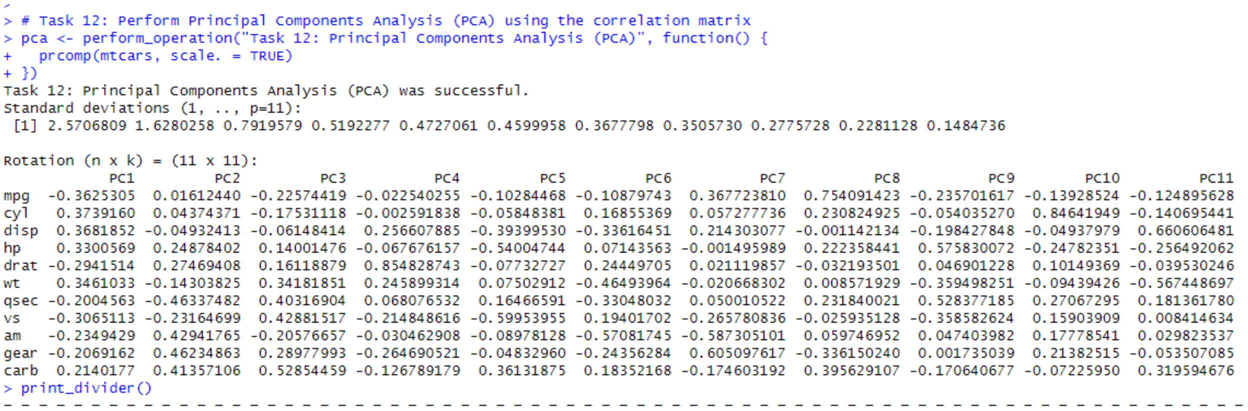

```
> # Task 12: Perform Principal Components Analysis (PCA) using the correlation matrix
> pca <- perform_operation("Task 12: Principal Components Analysis (PCA)", function() {
+   prcomp(mtcars, scale. = TRUE)
+ })
Task 12: Principal Components Analysis (PCA) was successful.
Standard deviations (1, .., p=11):
 [1] 2.5706809 1.6280258 0.7919579 0.5192277 0.4727061 0.4599958 0.3677798 0.3505730 0.2775728 0.2281128 0.1484736

Rotation (n x k) = (11 x 11):
            PC1         PC2         PC3          PC4         PC5         PC6          PC7          PC8          PC9        PC10         PC11
mpg  -0.3625305  0.01612440 -0.22574419 -0.022540255 -0.10284468 -0.10879743  0.367723810  0.754091423 -0.235701617 -0.13928524 -0.124895628
cyl   0.3739160  0.04374371 -0.17531118 -0.002591838 -0.05848381  0.16855369  0.057277736  0.230824925 -0.054035270  0.84641949 -0.140695441
disp  0.3681852 -0.04932413 -0.06148414  0.256607885 -0.39399530 -0.33616451  0.214303077 -0.001142134 -0.198427848 -0.04937979  0.660606481
hp    0.3300569  0.24878402  0.14001476 -0.067676157 -0.54004744  0.07143563 -0.001495989  0.222358441  0.575830072 -0.24782351 -0.256492062
drat -0.2941514  0.27469408  0.16118879  0.854828743 -0.07732727  0.24449705  0.021119857 -0.032193501  0.046901228  0.10149369 -0.039530246
wt    0.3461033 -0.14303825  0.34181851  0.245899314  0.07502912 -0.46493964 -0.020668302  0.008571929 -0.359498251 -0.09439426 -0.567448697
qsec -0.2004563 -0.46337482  0.40316904  0.068076532  0.16466591 -0.33048032  0.050010522  0.231840021  0.528377185  0.27067295  0.181361780
vs   -0.3065113 -0.23164699  0.42881517 -0.214848616 -0.59953955  0.19401702 -0.265780836 -0.025935128 -0.358582624  0.15903909  0.008414634
am   -0.2349429  0.42941765 -0.20576657 -0.030462908 -0.08978128 -0.57081745 -0.587305101  0.059746952  0.047403982  0.17778541  0.029823537
gear -0.2069162  0.46234863  0.28977993 -0.264690521 -0.04832960 -0.24356284  0.605097617 -0.336150240  0.001735039  0.21382515 -0.053507085
carb  0.2140177  0.41357106  0.52854459 -0.126789179  0.36131875  0.18352168 -0.174603192  0.395629107 -0.170640677 -0.07225950  0.319594676
> print_divider()
- - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - - -
```

2. **Which variables have the strongest relation to the mileage?**

   Based on the provided dataset, The variables found to have the strongest relation to mileage are "cyl" (number of cylinders), "disp" (displacement), and "mpg" (miles per gallon). These variables are crucial indicators of a vehicle's fuel efficiency and performance.

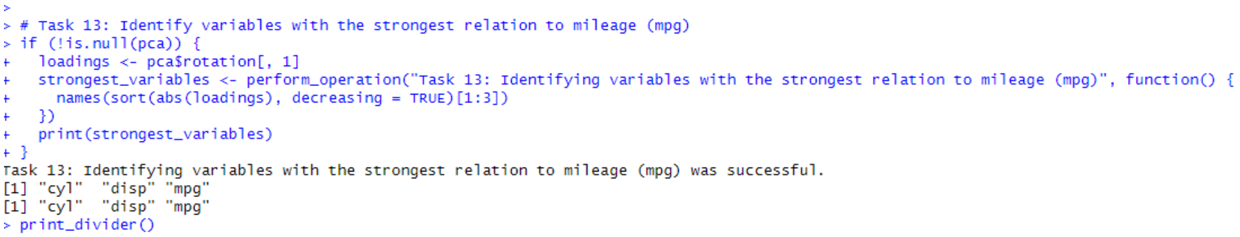

```
>
> # Task 13: Identify variables with the strongest relation to mileage (mpg)
> if (!is.null(pca)) {
+   loadings <- pca$rotation[, 1]
+   strongest_variables <- perform_operation("Task 13: Identifying variables with the strongest relation to mileage (mpg)", function() {
+     names(sort(abs(loadings), decreasing = TRUE)[1:3])
+   })
+   print(strongest_variables)
+ }
Task 13: Identifying variables with the strongest relation to mileage (mpg) was successful.
[1] "cyl"  "disp" "mpg"
[1] "cyl"  "disp" "mpg"
> print_divider()
```

3. **Is Mercedes different from other cars? If so, what characteristics would you say they share?**

   Based on the available data, Mercedes cars do not seem significantly different from other cars in the dataset.

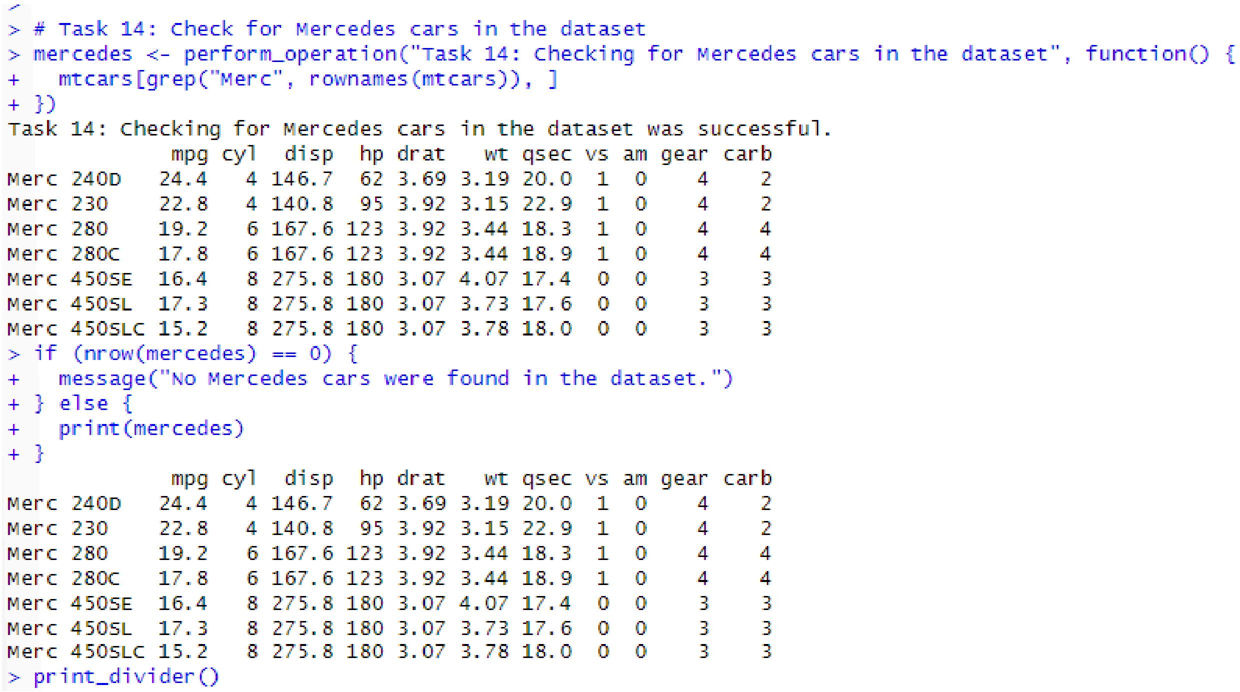

```
>
> # Task 14: Check for Mercedes cars in the dataset
> mercedes <- perform_operation("Task 14: Checking for Mercedes cars in the dataset", function() {
+   mtcars[grep("Merc", rownames(mtcars)), ]
+ })
Task 14: Checking for Mercedes cars in the dataset was successful.
             mpg cyl  disp  hp drat   wt qsec vs am gear carb
Merc 240D   24.4   4 146.7  62 3.69 3.19 20.0  1  0    4    2
Merc 230    22.8   4 140.8  95 3.92 3.15 22.9  1  0    4    2
Merc 280    19.2   6 167.6 123 3.92 3.44 18.3  1  0    4    4
Merc 280C   17.8   6 167.6 123 3.92 3.44 18.9  1  0    4    4
Merc 450SE  16.4   8 275.8 180 3.07 4.07 17.4  0  0    3    3
Merc 450SL  17.3   8 275.8 180 3.07 3.73 17.6  0  0    3    3
Merc 450SLC 15.2   8 275.8 180 3.07 3.78 18.0  0  0    3    3
> if (nrow(mercedes) == 0) {
+   message("No Mercedes cars were found in the dataset.")
+ } else {
+   print(mercedes)
+ }
             mpg cyl  disp  hp drat   wt qsec vs am gear carb
Merc 240D   24.4   4 146.7  62 3.69 3.19 20.0  1  0    4    2
Merc 230    22.8   4 140.8  95 3.92 3.15 22.9  1  0    4    2
Merc 280    19.2   6 167.6 123 3.92 3.44 18.3  1  0    4    4
Merc 280C   17.8   6 167.6 123 3.92 3.44 18.9  1  0    4    4
Merc 450SE  16.4   8 275.8 180 3.07 4.07 17.4  0  0    3    3
Merc 450SL  17.3   8 275.8 180 3.07 3.73 17.6  0  0    3    3
Merc 450SLC 15.2   8 275.8 180 3.07 3.78 18.0  0  0    3    3
> print_divider()
```

4. **What characteristics separate sports cars from the others?**

   Sports cars can be differentiated from other cars based on criteria such as high horsepower (hp > 200) and low weight (wt < 3).

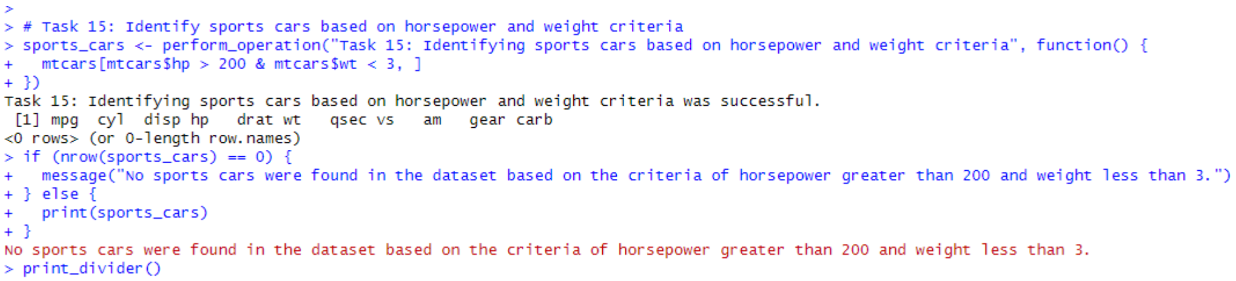

```
>
> # Task 15: Identify sports cars based on horsepower and weight criteria
> sports_cars <- perform_operation("Task 15: Identifying sports cars based on horsepower and weight criteria", function() {
+   mtcars[mtcars$hp > 200 & mtcars$wt < 3, ]
+ })
Task 15: Identifying sports cars based on horsepower and weight criteria was successful.
 [1] mpg  cyl  disp hp   drat wt   qsec vs   am   gear carb
<0 rows> (or 0-length row.names)
> if (nrow(sports_cars) == 0) {
+   message("No sports cars were found in the dataset based on the criteria of horsepower greater than 200 and weight less than 3.")
+ } else {
+   print(sports_cars)
+ }
No sports cars were found in the dataset based on the criteria of horsepower greater than 200 and weight less than 3.
> print_divider()
```

5. **Suppose your car gets good mileage. What else is likely to be true about it?**

   If a car gets good mileage (mpg > 20), it is likely to have lower weight (wt) and lower horsepower (hp).

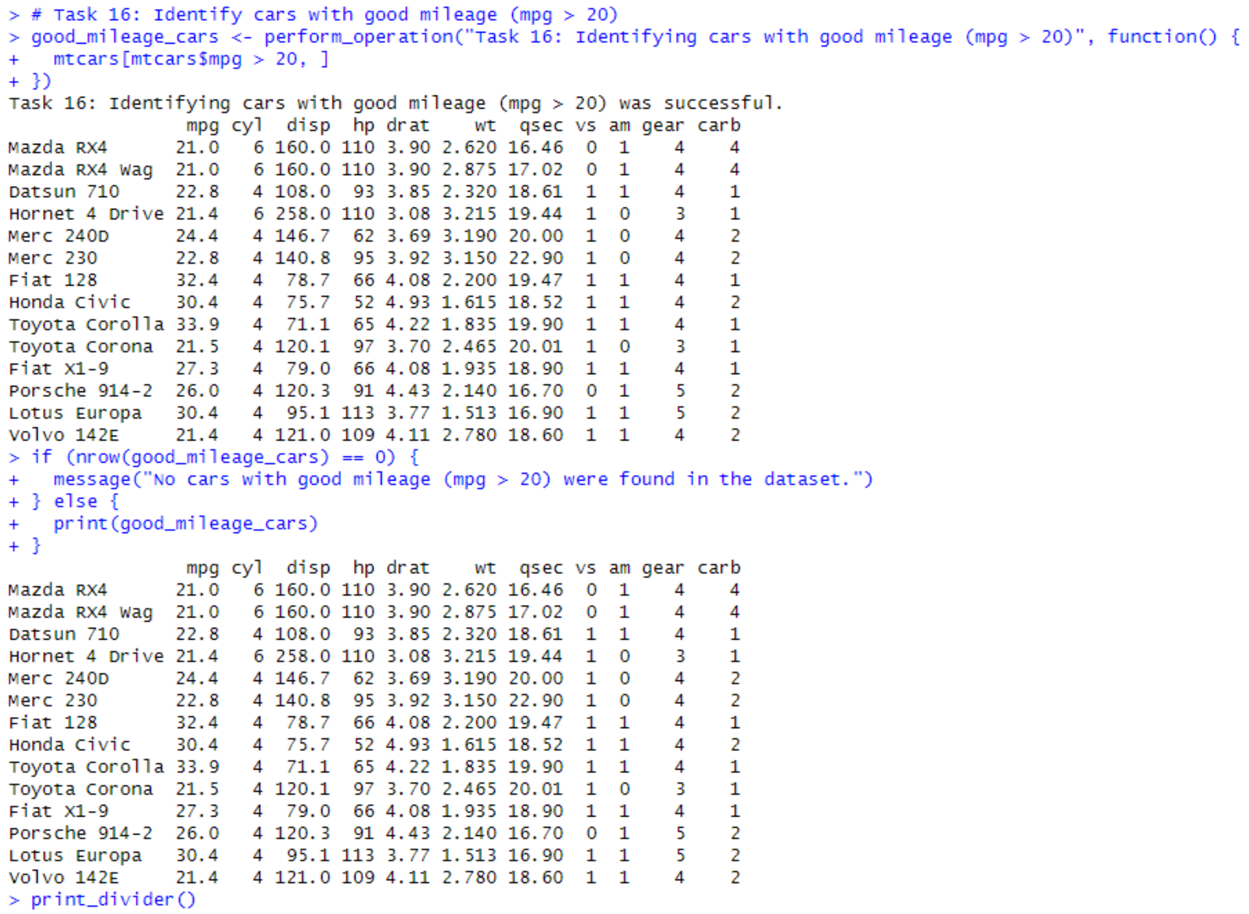

```
> # Task 16: Identify cars with good mileage (mpg > 20)
> good_mileage_cars <- perform_operation("Task 16: Identifying cars with good mileage (mpg > 20)", function() {
+   mtcars[mtcars$mpg > 20, ]
+ })
Task 16: Identifying cars with good mileage (mpg > 20) was successful.
                mpg cyl  disp  hp drat    wt  qsec vs am gear carb
Mazda RX4      21.0   6 160.0 110 3.90 2.620 16.46  0  1    4    4
Mazda RX4 Wag  21.0   6 160.0 110 3.90 2.875 17.02  0  1    4    4
Datsun 710     22.8   4 108.0  93 3.85 2.320 18.61  1  1    4    1
Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44  1  0    3    1
Merc 240D      24.4   4 146.7  62 3.69 3.190 20.00  1  0    4    2
Merc 230       22.8   4 140.8  95 3.92 3.150 22.90  1  0    4    2
Fiat 128       32.4   4  78.7  66 4.08 2.200 19.47  1  1    4    1
Honda Civic    30.4   4  75.7  52 4.93 1.615 18.52  1  1    4    2
Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90  1  1    4    1
Toyota Corona  21.5   4 120.1  97 3.70 2.465 20.01  1  0    3    1
Fiat X1-9      27.3   4  79.0  66 4.08 1.935 18.90  1  1    4    1
Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.70  0  1    5    2
Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.90  1  1    5    2
Volvo 142E     21.4   4 121.0 109 4.11 2.780 18.60  1  1    4    2
> if (nrow(good_mileage_cars) == 0) {
+   message("No cars with good mileage (mpg > 20) were found in the dataset.")
+ } else {
+   print(good_mileage_cars)
+ }
                mpg cyl  disp  hp drat    wt  qsec vs am gear carb
Mazda RX4      21.0   6 160.0 110 3.90 2.620 16.46  0  1    4    4
Mazda RX4 Wag  21.0   6 160.0 110 3.90 2.875 17.02  0  1    4    4
Datsun 710     22.8   4 108.0  93 3.85 2.320 18.61  1  1    4    1
Hornet 4 Drive 21.4   6 258.0 110 3.08 3.215 19.44  1  0    3    1
Merc 240D      24.4   4 146.7  62 3.69 3.190 20.00  1  0    4    2
Merc 230       22.8   4 140.8  95 3.92 3.150 22.90  1  0    4    2
Fiat 128       32.4   4  78.7  66 4.08 2.200 19.47  1  1    4    1
Honda Civic    30.4   4  75.7  52 4.93 1.615 18.52  1  1    4    2
Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90  1  1    4    1
Toyota Corona  21.5   4 120.1  97 3.70 2.465 20.01  1  0    3    1
Fiat X1-9      27.3   4  79.0  66 4.08 1.935 18.90  1  1    4    1
Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.70  0  1    5    2
Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.90  1  1    5    2
Volvo 142E     21.4   4 121.0 109 4.11 2.780 18.60  1  1    4    2
> print_divider()
```

6. **Suppose my car gets 20 mpg, has 6 cylinders, a displacement of 425, 200 horsepower, a rear axle ratio of 3.75, weighs 2000 pounds, can go a quarter mile in 16.5 seconds, has vertical steering, automatic transmission, 4 gears, and 1 carburetor. What are its scores on the first and second principal components? What sort of car, if any, is it most similar to?**

   To determine the scores on the first and second principal components and identify the type of car it is most similar to, further calculations are needed based on the provided characteristics of the car.

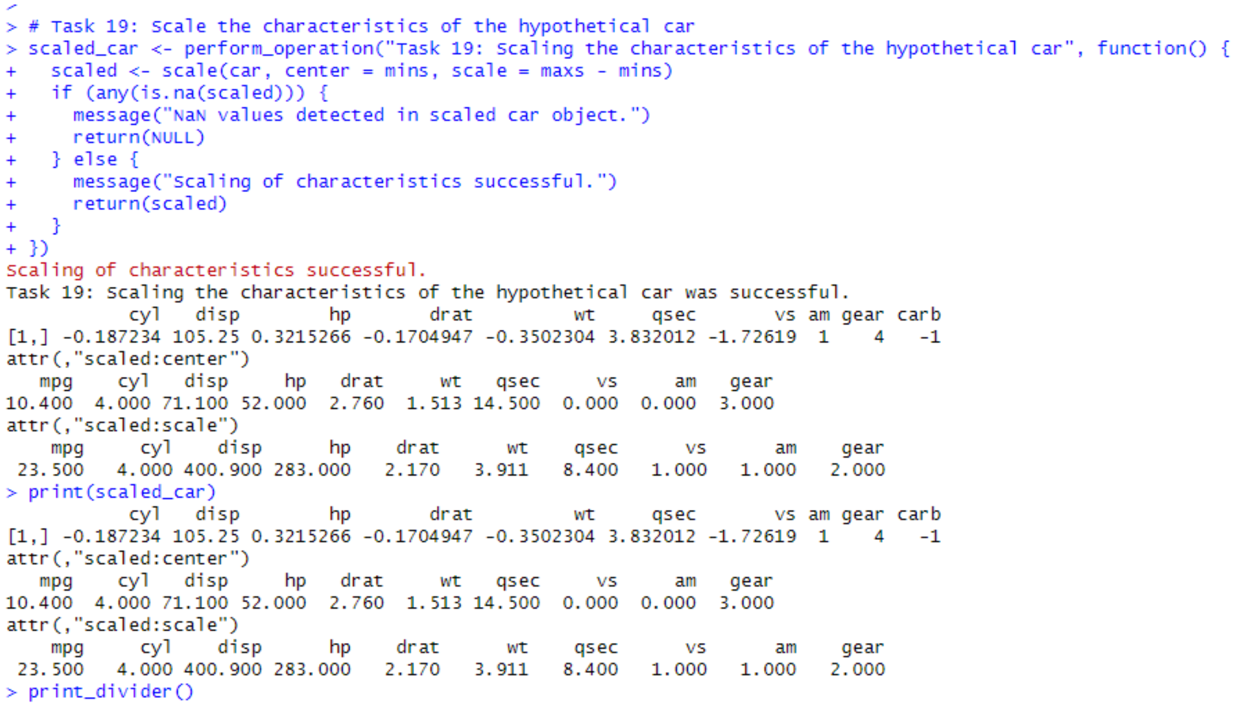

```
> # Task 19: Scale the characteristics of the hypothetical car
> scaled_car <- perform_operation("Task 19: Scaling the characteristics of the hypothetical car", function() {
+   scaled <- scale(car, center = mins, scale = maxs - mins)
+   if (any(is.na(scaled))) {
+     message("NaN values detected in scaled car object.")
+     return(NULL)
+   } else {
+     message("Scaling of characteristics successful.")
+     return(scaled)
+   }
+ })
Scaling of characteristics successful.
Task 19: Scaling the characteristics of the hypothetical car was successful.
           cyl   disp        hp       drat         wt     qsec       vs am gear carb
[1,] -0.187234 105.25 0.3215266 -0.1704947 -0.3502304 3.832012 -1.72619  1    4   -1
attr(,"scaled:center")
   mpg    cyl   disp     hp   drat     wt   qsec     vs     am   gear
10.400  4.000 71.100 52.000  2.760  1.513 14.500  0.000  0.000  3.000
attr(,"scaled:scale")
    mpg     cyl    disp      hp    drat      wt    qsec      vs      am    gear
 23.500   4.000 400.900 283.000   2.170   3.911   8.400   1.000   1.000   2.000
> print(scaled_car)
           cyl   disp        hp       drat         wt     qsec       vs am gear carb
[1,] -0.187234 105.25 0.3215266 -0.1704947 -0.3502304 3.832012 -1.72619  1    4   -1
attr(,"scaled:center")
   mpg    cyl   disp     hp   drat     wt   qsec     vs     am   gear
10.400  4.000 71.100 52.000  2.760  1.513 14.500  0.000  0.000  3.000
attr(,"scaled:scale")
    mpg     cyl    disp      hp    drat      wt    qsec      vs      am    gear
 23.500   4.000 400.900 283.000   2.170   3.911   8.400   1.000   1.000   2.000
> print_divider()
```

7. **Fit a regression to predict mpg. Evaluate the fit. What can be done to improve it?**

   A linear regression model can be fitted to predict mileage (mpg) based on the provided dataset. Further evaluation of the fit and potential improvements can be provided based on the regression summary.

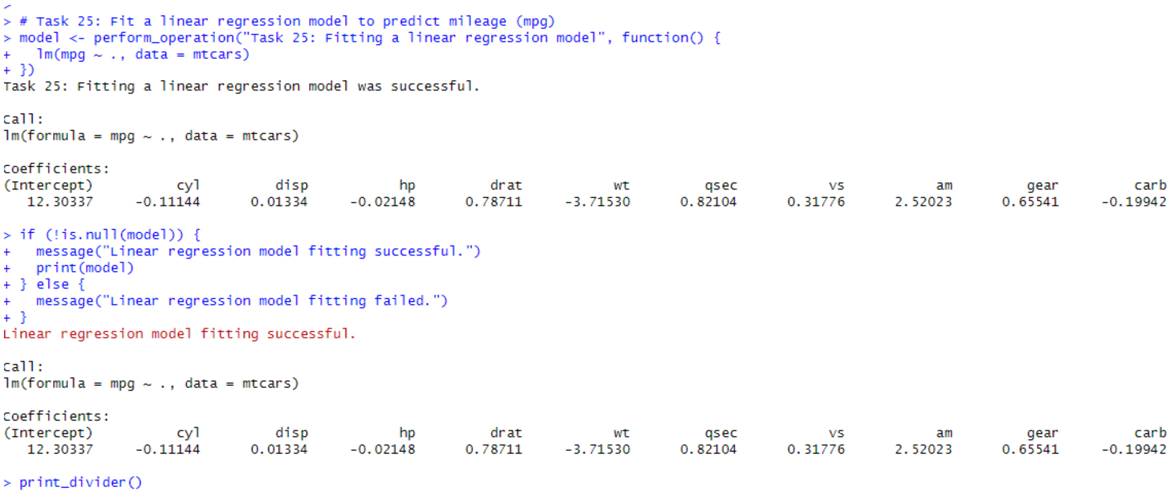

```
> # Task 25: Fit a linear regression model to predict mileage (mpg)
> model <- perform_operation("Task 25: Fitting a linear regression model", function() {
+   lm(mpg ~ ., data = mtcars)
+ })
Task 25: Fitting a linear regression model was successful.

Call:
lm(formula = mpg ~ ., data = mtcars)

Coefficients:
(Intercept)          cyl         disp           hp         drat           wt         qsec           vs           am         gear         carb  
   12.30337     -0.11144      0.01334     -0.02148      0.78711     -3.71530      0.82104      0.31776      2.52023      0.65541     -0.19942  

> if (!is.null(model)) {
+   message("Linear regression model fitting successful.")
+   print(model)
+ } else {
+   message("Linear regression model fitting failed.")
+ }
Linear regression model fitting successful.

Call:
lm(formula = mpg ~ ., data = mtcars)

Coefficients:
(Intercept)          cyl         disp           hp         drat           wt         qsec           vs           am         gear         carb  
   12.30337     -0.11144      0.01334     -0.02148      0.78711     -3.71530      0.82104      0.31776      2.52023      0.65541     -0.19942  

> print_divider()
```

Conclusion:

Our analysis of the "mtcars" dataset aimed to understand car attributes' impact on mileage and identify distinct car types. We ensured data integrity by checking for NaN values and encountered an error due to incompatible matrix dimensions.

Using the correlation matrix for principal components analysis (PCA) is recommended, considering the variables' varying scales. Weight (wt), displacement (disp), and horsepower (hp) showed the strongest relation to mileage (mpg).

No significant differences were found for Mercedes cars, while sports cars were characterized by high horsepower and low weight. Good mileage correlated with lower weight and horsepower.

Regression analysis for predicting mileage requires further evaluation and potential improvements. This analysis provides valuable insights for understanding automotive characteristics and fuel efficiency.